

Topics in Statistical Theory — Summary

Lucas Riedstra

November 21, 2020

Contents

1	Basic concepts	2
1.1	Parametric vs nonparametric models	2
1.2	Estimating an arbitrary distribution function	2
1.3	Order statistics and quantiles	5
1.4	Concentration inequalities	6
2	Kernel density estimation	12
2.1	The univariate kernel density estimator	12
2.2	Bounds on variance and bias	13
2.3	Bounds on the integrated variance and bias	15
2.4	Bandwidth selection	16
2.4.1	Least squares cross validation	16
2.4.2	Lepski	17
2.5	Choice of kernel	17
2.6	Multivariate density estimation	17
3	Nonparametric regression	18
3.1	Fixed and random design	18
3.2	Local polynomial estimators	18
3.3	Splines	22
3.3.1	Cubic splines	22
3.3.2	Natural cubic smoothing splines	23
3.3.3	Choice of smoothing parameter	23
4	Minimax lower bounds	25
4.1	Reduction to testing	25
4.2	Divergences	26

1 Basic concepts

1.1 Parametric vs nonparametric models

Definition 1.1. A *statistical model* is a family of possible data-generating mechanisms. If the parameter space Θ is finite-dimensional, we speak of a *parametric model*.

A model is called *well-specified* if there is a $\vartheta_0 \in \Theta$ for which the data was generated from the distribution with parameter ϑ_0 , and otherwise it is called *misspecified*.

Recap 1.2. Let (Y_n) be a sequence of random vectors and Y a random vector.

1. We say that (Y_n) *converges almost surely* to Y , notation $Y_n \xrightarrow{\text{a.s.}} Y$, if $\mathbb{P}(Y_n \rightarrow Y) = 1$.
2. We say that (Y_n) *converges in probability* to Y , notation $Y_n \xrightarrow{\text{P}} Y$, if for every $\varepsilon > 0$ we have $\mathbb{P}(\|Y_n - Y\| > \varepsilon) \rightarrow 0$.
3. We say that (Y_n) *converges in distribution* to Y , notation $Y_n \xrightarrow{\text{d}} Y$, if $\mathbb{P}(Y_n \leq y) \rightarrow \mathbb{P}(Y \leq y)$ for all y where the distribution function of Y is continuous.

This is equivalent to the condition that $\mathbb{E}[f(Y_n)] \rightarrow \mathbb{E}[f(Y)]$ for all bounded Lipschitz functions f .

It is known that $Y_n \xrightarrow{\text{a.s.}} Y \implies Y_n \xrightarrow{\text{P}} Y \implies Y_n \xrightarrow{\text{d}} Y$.

If (Y_n) is a sequence of random vectors and (a_n) is a positive sequence, then we write $Y_n = O_p(a_n)$ if, for all $\varepsilon > 0$, there exists $C > 0$ such that for sufficiently large n we have

$$\mathbb{P}\left(\frac{\|Y_n\|}{a_n} > C\right) < \varepsilon.$$

We write $Y_n = o_p(a_n)$ if $Y_n/a_n \xrightarrow{\text{P}} 0$.

In a well-specified parametric model, the maximum likelihood estimator (MLE) $\hat{\vartheta}_n$ typically satisfies $\hat{\vartheta}_n - \vartheta_0 \in O_p(n^{-1/2})$. On the other hand, if the model is misspecified, any inference can give very misleading results. To circumvent this problem, we consider *nonparametric models*, which make much weaker assumptions. Such infinite-dimensional models are much less vulnerable to model misspecification, however we will typically pay a price in terms of a slower convergence rate than in well-specified parametric models.

Example 1.3. Examples of nonparametric models include:

1. Assume $X_1, \dots, X_n \stackrel{\text{iid}}{\sim} F$ for some unknown distribution function F .
2. Assume $X_1, \dots, X_n \stackrel{\text{iid}}{\sim} f$ for some unknown density f belonging to a smoothness class.
3. Assume $Y_i = m(x_i) + \varepsilon_i$ ($i = 1, \dots, n$), where the x_i are known, m is unknown and belongs to some smoothness class, and the ε_i are i.i.d. with $\mathbb{E}(\varepsilon_i) = 0$ and $\text{Var}(\varepsilon_i) = \sigma^2$.

1.2 Estimating an arbitrary distribution function

Definition 1.4. Let \mathcal{F} denote the class of all distribution functions on \mathbb{R} and suppose $X_1, \dots, X_n \stackrel{\text{iid}}{\sim} F \in \mathcal{F}$. The *empirical distribution function* \hat{F}_n of X_1, \dots, X_n is defined as

$$\hat{F}_n(x) := \frac{1}{n} \sum_{i=1}^n \mathbb{1}_{\{X_i \leq x\}}.$$

Recap 1.5. The *strong law of large numbers* tells us that if (Y_n) are i.i.d. with finite mean μ , then $\bar{Y} := \frac{1}{n} \sum_{i=1}^n Y_i \xrightarrow{\text{a.s.}} \mu$.

Note that the strong law of large numbers immediately implies that $\hat{F}_n(x)$ converges almost surely to $F(x)$ as $n \rightarrow \infty$. However, the following stronger result states that this convergence holds uniformly in x :

Theorem 1.6 (Glivenko-Cantelli). *Let $X_1, X_2, \dots \stackrel{\text{iid}}{\sim} F$. Then we have*

$$\sup_{x \in \mathbb{R}} \left| \hat{F}_n(x) - F(x) \right| \xrightarrow{\text{a.s.}} 0.$$

Proof. See lecture notes. The main idea of the proof is to “control” \hat{F}_n in a finite number of points x_1, \dots, x_k , and then deduce what happens between those points using the fact that distributions are increasing and right-continuous. On [Wikipedia](#), a simplified proof can be found assuming that F is continuous, which still encapsulates the main idea. \square

Theorem 1.7 (Dvoretzky-Kiefer-Wolfowitz). *Under the conditions of theorem 1.6, for every $\varepsilon > 0$ it holds that*

$$\mathbb{P}_F \left(\sup_{x \in \mathbb{R}} \left| \hat{F}_n(x) - F(x) \right| > \varepsilon \right) \leq 2e^{-2n\varepsilon^2},$$

and this is a tight bound.

We will not prove this theorem, however, we will explore a few consequences. One of these consequences is the following:

Corollary 1.8 (Uniform Glivenko-Cantelli theorem). *Under the conditions of theorem 1.6, for every $\varepsilon > 0$, it holds that*

$$\sup_{F \in \mathcal{F}} \mathbb{P}_F \left(\sup_{m \geq n} \sup_{x \in \mathbb{R}} \left| \hat{F}_m(x) - F(x) \right| > \varepsilon \right) \rightarrow 0 \quad \text{as } n \rightarrow \infty.$$

Proof. By a union bound, the DKW inequality, and convergence of the geometric series we have

$$\begin{aligned} \sup_{F \in \mathcal{F}} \mathbb{P}_F \left(\sup_{m \geq n} \sup_{x \in \mathbb{R}} \left| \hat{F}_m(x) - F(x) \right| > \varepsilon \right) &\leq \sup_{F \in \mathcal{F}} \sum_{m=n}^{\infty} \mathbb{P}_F \left(\sup_{x \in \mathbb{R}} \left| \hat{F}_m(x) - F(x) \right| > \varepsilon \right) \\ &\leq 2 \sum_{m=n}^{\infty} e^{-2m\varepsilon^2}, \end{aligned}$$

which converges to 0 as it is the tail of a converging sum. \square

For another consequence, we consider the problem of finding a confidence band for F . Given $\alpha \in (0, 1)$, set $\varepsilon_n := \sqrt{-\frac{1}{2n} \log(\alpha/2)}$. Then the DKW inequality tells us that

$$\mathbb{P}_F \left(\sup_{x \in \mathbb{R}} \left| \hat{F}_n(x) - F(x) \right| > \varepsilon_n \right) \leq \alpha,$$

or equivalently, that

$$\mathbb{P}_F \left(\hat{F}_n(x) - \varepsilon_n \leq F(x) \leq \hat{F}_n(x) + \varepsilon_n \text{ for all } x \in \mathbb{R} \right) \geq 1 - \alpha.$$

We can say even more.

Recap 1.9. For any distribution function F , its *quantile function* is defined as

$$F^{-1}: (0, 1] \rightarrow \mathbb{R} \cup \{\infty\}: p \mapsto \inf \{x \in \mathbb{R} \mid F(x) \geq p\}.$$

When necessary, we also define $F^{-1}(0) := \sup \{x \in \mathbb{R} \mid F(x) = 0\}$.

If $U \sim U(0, 1)$ and $X \sim F$, then for any $x \in \mathbb{R}$ we have

$$\mathbb{P}(F^{-1}(U) \leq x) = \mathbb{P}(U \leq F(x)) = F(x) = \mathbb{P}(X \leq x).$$

This can be written simply as $F^{-1}(U) \stackrel{d}{=} X$.

Let $U_1, \dots, U_n \stackrel{\text{iid}}{\sim} U(0, 1)$ with empirical distribution function \hat{G}_n , and let $X_1, \dots, X_n \stackrel{\text{iid}}{\sim} F$. Then, we have

$$\hat{G}_n(F(x)) = \frac{1}{n} \sum_{i=1}^n \mathbb{1}_{\{U_i \leq F(x)\}} \stackrel{d}{=} \frac{1}{n} \sum_{i=1}^n \mathbb{1}_{\{X_i \leq x\}} = \hat{F}_n(x),$$

where $\stackrel{d}{=}$ means equality in distribution. It follows that

$$\sup_{x \in \mathbb{R}} \left| \hat{F}_n(x) - F(x) \right| \stackrel{d}{=} \sup_{x \in \mathbb{R}} \left| \hat{G}_n(F(x)) - F(x) \right| \leq \sup_{t \in [0, 1]} \left| \hat{G}_n(t) - t \right|,$$

with equality if F is continuous. We conclude that if F is continuous, the distribution of $\sup_{x \in \mathbb{R}} \left| \hat{F}_n(x) - F(x) \right|$ does not depend on F .

Other generalisations of theorem 1.6 include Uniform Laws of Large Numbers. Let X, X_1, \dots, X_n be i.i.d. on a measurable space $(\mathcal{X}, \mathcal{A})$, and \mathcal{G} a class of measurable functions on \mathcal{X} . We say that \mathcal{G} satisfies a ULLN if

$$\sup_{g \in \mathcal{G}} \left| \frac{1}{n} \sum_{i=1}^n g(X_i) - \mathbb{E}[g(X)] \right| \xrightarrow{\text{a.s.}} 0.$$

In theorem 1.6, we showed that $\mathcal{G} = \{\mathbb{1}_{\{\cdot \leq x\}} \mid x \in \mathbb{R}\}$ satisfies a ULLN.

Recap 1.10. We recall the central limit theorem: if X_1, \dots, X_n are i.i.d. random variables with mean μ and variance $\sigma^2 < \infty$, then $\sqrt{n}(\bar{X}_n - \mu) \xrightarrow{d} N(0, \sigma^2)$.

Dividing by σ yields

$$\frac{\sqrt{n}(\bar{X}_n - \mu)}{\sigma} \xrightarrow{d} N(0, 1),$$

and multiplying both sides by n and writing $V_i = \sum_{j=1}^i X_j$ we obtain

$$\frac{V_i - \mathbb{E}V_i}{\sqrt{\text{Var}(V_i)}} \xrightarrow{d} N(0, 1).$$

Another extension starts with the observation that $\sqrt{n}(\hat{F}_n(x) - F(x)) \xrightarrow{d} N(0, \sigma^2)$, where

$$\sigma^2 = \text{Var}(\mathbb{1}_{\{X \leq x\}}) = \mathbb{E}[\mathbb{1}_{\{X \leq x\}}^2] - \mathbb{E}[\mathbb{1}_{\{X \leq x\}}]^2 = F(x) - F(x)^2 = F(x)(1 - F(x)).$$

This can be strengthened by considering $(\sqrt{n}(\hat{F}_n(x) - F(x)) \xrightarrow{d} N(0, \sigma^2) \mid x \in \mathbb{R})$ as a stochastic process.

1.3 Order statistics and quantiles

Definition 1.11. Let $X_1, \dots, X_n \stackrel{\text{iid}}{\sim} F \in \mathcal{F}$. The *order statistics* are the ordered samples $X_{(1)} \leq \dots \leq X_{(n)}$ (where the original order is preserved in case of a tie).

The order statistics of the uniform distribution can be computed explicitly:

Proposition 1.12. Let $U_1, \dots, U_n \stackrel{\text{iid}}{\sim} U(0, 1)$, let $Y_1, \dots, Y_{n+1} \stackrel{\text{iid}}{\sim} \text{Exp}(1)$, and write $S_j := \sum_{i=1}^j Y_i$ ($j = 1, \dots, n+1$). Then

$$U_{(j)} \stackrel{d}{=} \frac{S_j}{S_{n+1}} \sim \text{Beta}(j, n-j+1) \quad \text{for } j = 1, \dots, n.$$

Proof. See example sheet 1, question 1. □

Definition 1.13. Let $X_1, \dots, X_n \stackrel{\text{iid}}{\sim} F$. Then the *sample quantile function* is defined as

$$\hat{F}_n^{-1}(p) = \inf \left\{ x \in \mathbb{R} \mid \hat{F}_n(x) \geq p \right\}.$$

Proposition 1.14. It holds that $\hat{F}_n^{-1}(p) = X_{(\lceil np \rceil)}$.

Proof. By definition, $\hat{F}_n^{-1}(p)$ is the smallest value of x for which $\hat{F}_n(x)$ is larger than p . Note that

$$\hat{F}_n(x) \geq p \iff \frac{1}{n} \sum_{i=1}^n \mathbb{1}_{\{X_i \leq x\}} \geq p \iff \sum_{i=1}^n \mathbb{1}_{\{X_i \leq x\}} \geq np \iff \sum_{i=1}^n \mathbb{1}_{\{X_i \leq x\}} \geq \lceil np \rceil.$$

The smallest value of x for which this occurs is the smallest value of x such that exactly $\lceil np \rceil$ of the variables X_1, \dots, X_n satisfy $X_i \leq x$. We conclude that $\hat{F}_n^{-1}(p) = X_{(\lceil np \rceil)}$. □

For $p = \frac{1}{2}$ for example, this proposition tells us that $\hat{F}_n^{-1}(p) = X_{(\lceil n/2 \rceil)}$, the median of the data. We now explore the distribution of $X_{(\lceil np \rceil)}$.

Recap 1.15. We recall two theorems. The first is *Slutsky's theorem*:

Theorem 1.16. Let (Y_n) and (Z_n) be sequences of random vectors with $Y_n \xrightarrow{d} Y$ and $Z_n \xrightarrow{P} c$ for some constant c . If g is a continuous real-valued function, then $g(Y_n, Z_n) \xrightarrow{d} g(Y, c)$.

The second is the *delta method*:

Theorem 1.17. Let (Y_n) be a sequence of random vectors such that $\sqrt{n}(Y_n - \mu) \xrightarrow{d} Z$. If $g: \mathbb{R}^d \rightarrow \mathbb{R}$ is differentiable at μ , then

$$\sqrt{n}(g(Y_n) - g(\mu)) \xrightarrow{d} g'(\mu)Z.$$

Lemma 1.18. If $U_1, \dots, U_n \stackrel{\text{iid}}{\sim} U(0, 1)$ and $p \in (0, 1)$, then $\sqrt{n}(U_{\lceil np \rceil} - p) \xrightarrow{d} N(0, p(1-p))$.

Proof. Let $Y_1, \dots, Y_{n+1} \stackrel{\text{iid}}{\sim} \text{Exp}(1)$, $V_n := \sum_{i=1}^{\lceil np \rceil} Y_i$ and $W_n := \sum_{i=\lceil np \rceil+1}^{n+1} Y_i$. Then V_n and W_n are independent, and we have seen that $U_{\lceil np \rceil} \sim \frac{V_n}{V_n + W_n}$.

Noting that $\mathbb{E}V_n = \text{Var}(V_n) = \lceil np \rceil$ we find

$$\begin{aligned} \sqrt{n} \left(\frac{V_n}{n} - p \right) &= \frac{\sqrt{\lceil np \rceil}}{\sqrt{n}} \left(\frac{V_n - \lceil np \rceil}{\sqrt{\lceil np \rceil}} \right) + \frac{\lceil np \rceil - np}{\sqrt{n}} \\ &= \frac{\sqrt{\lceil np \rceil}}{\sqrt{n}} \left(\frac{V_n - \mathbb{E}V_n}{\sqrt{\text{Var}(V_n)}} \right) + \frac{\lceil np \rceil - np}{\sqrt{n}}. \end{aligned}$$

Now, by the central limit theorem, the term between brackets converges to a standard $N(0, 1)$ distribution. The term $\sqrt{\lceil np \rceil} \sqrt{n}$ converges to \sqrt{p} and the term $(\lceil np \rceil - np)/\sqrt{n}$ converges to 0, so by Slutsky's lemma, we find

$$\sqrt{n} \left(\frac{V_n}{n} - p \right) \xrightarrow{d} \sqrt{p} N(0, 1) = N(0, p).$$

An analogous calculation shows that $\sqrt{n} \left(\frac{W_n}{n} - (1 - p) \right) \rightarrow N(0, 1 - p)$.

Now we define $g: (0, \infty)^2 \rightarrow (0, \infty)$ by $g(x, y) := x/(x + y)$, which is differentiable at $(p, 1 - p)$. Note that the distribution of (V_n, W_n) is an $N(0, \begin{pmatrix} p & 0 \\ 0 & q \end{pmatrix})$ distribution. By the delta method we find

$$\begin{aligned} \sqrt{n}(U_{\lceil np \rceil} - p) &\stackrel{d}{=} \sqrt{n} \left(g \left(\frac{V_n}{n}, \frac{W_n}{n} \right) - g(p, q) \right) \\ &\stackrel{d}{\rightarrow} g'(p, 1 - p) N \left(0, \begin{pmatrix} p & 0 \\ 0 & q \end{pmatrix} \right) \\ &= N \left(0, g'(p, 1 - p) \begin{pmatrix} p & 0 \\ 0 & q \end{pmatrix} g'(p, 1 - p)^\top \right) \\ &= N(0, p(1 - p)). \end{aligned}$$

□

We now relate what we know about the uniform distribution to the quantile function:

Theorem 1.19. *Let $p \in (0, 1)$ and let $X_1, \dots, X_n \stackrel{\text{iid}}{\sim} F$. Suppose that F is differentiable at $\xi_p := F^{-1}(p)$ with derivative $f(\xi_p)$. Then*

$$\sqrt{n}(X_{(\lceil np \rceil)} - \xi_p) \xrightarrow{d} N \left(0, \frac{p(1 - p)}{f(\xi_p)^2} \right).$$

Proof. Let $U_1, \dots, U_n \stackrel{\text{iid}}{\sim} U(0, 1)$, then we know that $F^{-1}(U_i) \stackrel{d}{=} X_i$ and thus $F^{-1}(U_{(\lceil np \rceil)}) \stackrel{d}{=} X_{(\lceil np \rceil)}$. Applying the delta method with $g = F^{-1}$, together with the previous theorem yields

$$\sqrt{n}(X_{(\lceil np \rceil)} - \xi_p) = \sqrt{n}(F^{-1}(U_{(\lceil np \rceil)}) - F^{-1}(p)) \xrightarrow{d} (F^{-1})'(p) \cdot N(0, p(1 - p)).$$

Noting that $(F^{-1})'(p) = \frac{1}{f(\xi_p)}$ yields the result. □

1.4 Concentration inequalities

We turn our attention to concentration inequalities, with a focus on finite-sample results (instead of results that only hold for $n \rightarrow \infty$).

Definition 1.20. A random variable X with mean 0 is called *sub-Gaussian* with parameter σ^2 if

$$M_X(t) = \mathbb{E}(e^{tX}) \leq e^{t^2 \sigma^2 / 2}$$

for every $t \in \mathbb{R}$.

Note that equality holds when $X \sim N(0, \sigma^2)$, since the MGF of an $N(\mu, \sigma^2)$ distribution is given by $t \mapsto \exp(\mu t + \sigma^2 t^2 / 2)$.

Recap 1.21. Recall the *tail bound formula* for the expectation: if X is a nonnegative random variable, then

$$\mathbb{E}[X] = \int_0^\infty \mathbb{P}(X > x) dx.$$

Furthermore, recall that the *gamma function* is defined for $z \in (0, \infty)$ by

$$\Gamma(z) = \int_0^\infty x^{z-1} e^{-x} dx$$

and satisfies $\Gamma(n) = (n-1)!$ for all $n \in \mathbb{N}$.

Finally, recall the following inequality: for all $a, b \in \mathbb{R}$ and $p \geq 1$

$$(a+b)^p \leq 2^{p-1}(a^p + b^p).$$

This follows from the convexity of the function $x \mapsto x^p$.

Proposition 1.22. *We consider some characterisations of sub-Gaussianity:*

(a) *Let X be sub-Gaussian with parameter σ^2 . Then*

$$\max\{\mathbb{P}(X \geq x), \mathbb{P}(X \leq -x)\} \leq e^{-x^2/(2\sigma^2)} \quad \text{for every } x \geq 0. \quad (1)$$

(b) *Let X be a random variable which satisfies $\mathbb{E}(X) = 0$ and eq. (1). Then for every $q \in \mathbb{N}$ it holds that*

$$\mathbb{E}(X^{2q}) \leq 2 \cdot q!(2\sigma^2)^q \leq q!(2\sigma)^{2q}.$$

(c) *If X is a random variable with $\mathbb{E}(X) = 0$ and $\mathbb{E}(X^{2q}) \leq q!C^{2q}$ for all $q \in \mathbb{N}$, then X is sub-Gaussian with parameter $4C^2$.*

Proof. (a) We first consider $\mathbb{P}(X \geq x)$. By Markov's inequality, we have for all $t \in \mathbb{R}$ that

$$\mathbb{P}(X \geq x) = \mathbb{P}(e^{tX} \geq e^{tx}) \leq e^{-tx} \mathbb{E}(e^{tX}) \leq e^{-tx+t^2\sigma^2/2}.$$

Since the LHS is independent of t , we can take the infimum over t on the RHS and obtain

$$\mathbb{P}(X \geq x) \leq \inf_{t \in \mathbb{R}} e^{-tx+t^2\sigma^2/2} = e^{-x^2/(2\sigma^2)},$$

since the infimum of $t^2\sigma^2/2 - tx$ is attained at $t = x/\sigma^2$ (this method is called *Chernoff bounding*).

For $\mathbb{P}(X \leq -x) = \mathbb{P}(-X \geq x)$ we can use the fact that $-X$ is also sub-Gaussian with parameter σ^2 .

(b) By the previous part, we have $\mathbb{P}(|X| \geq x) \leq 2e^{-x^2/(2\sigma^2)}$. Some calculations give

$$\begin{aligned} \mathbb{E}(X^{2q}) &= \int_0^\infty \mathbb{P}(X^{2q} \geq x) dx = \int_0^\infty \mathbb{P}(|X| \geq x^{1/(2q)}) dx \\ &= 2q \int_0^\infty x^{2q-1} \mathbb{P}(|X| \geq x) dx \\ &\leq 4q \int_0^\infty x^{2q-1} e^{-x^2/(2\sigma^2)} dx. \end{aligned}$$

Now set $t = x^2/2\sigma^2$, so that $x = \sigma(2t)^{1/2}$ and thus $dx = \sigma(2t)^{-1/2} dt$. Plugging that in we get

$$\begin{aligned} \mathbb{E}(X^{2q}) &\leq 4q \int_0^\infty (\sigma(2t)^{1/2})^{2q-1} e^{-t} \sigma(2t)^{-1/2} dt = 2^{q+1} q \sigma^{2q} \int_0^\infty t^{q-1} e^{-t} dt \\ &= 2^{q+1} q \sigma^{2q} \Gamma(q) = 2 \cdot q!(2\sigma)^q. \end{aligned}$$

- (c) Note that $x \mapsto e^{-tx}$ is convex for every $t \in \mathbb{R}$, so $\mathbb{E}(e^{-tX}) \geq e^{-t\mathbb{E}(X)} = e^0 = 1$ by Jensen's inequality. Let X' denote an independent copy of X : then $X - X'$ has a symmetric distribution, so all its odd moments vanish. Therefore we find

$$\begin{aligned} \mathbb{E}[e^{tX}] &\leq \mathbb{E}[e^{-tX'}]\mathbb{E}[e^{tX}] = \mathbb{E}[e^{t(X-X')}] = \mathbb{E}\left[\sum_{q=0}^{\infty} \frac{t^{2q}(X-X')^{2q}}{(2q)!}\right] \\ &= \sum_{q=0}^{\infty} \frac{t^{2q}\mathbb{E}[(X-X')^{2q}]}{(2q)!} \leq \sum_{q=0}^{\infty} \frac{2^{2q-1}t^{2q}(\mathbb{E}[X^{2q}] + \mathbb{E}[(X')^{2q}])}{(2q)!} \\ &\leq \sum_{q=0}^{\infty} \frac{2^{2q-1}t^{2q}2q!C^{2q}}{(2q)!} = \sum_{q=0}^{\infty} \frac{(2tC)^{2q}q!}{(2q)!} = \sum_{q=0}^{\infty} \frac{(2tC)^{2q}}{\prod_{j=1}^q (q+j)} \\ &\leq \sum_{q=0}^{\infty} \frac{(2tC)^{2q}}{\prod_{j=1}^q (2j)} = \sum_{q=1}^{\infty} \frac{(2t^2C^2)^q}{q!} = e^{2t^2C^2}. \end{aligned}$$

This shows that X is sub-Gaussian with parameter $4C^2$. □

Note that the proposition is not an “if and only if”-type theorem: suppose we start with a sub-Gaussian variable X with parameter σ^2 . Then by (b), we have $\mathbb{E}[X^{2q}] \leq q!(2\sigma)^{2q}$, and (c) then implies that X is sub-Gaussian with parameter $16\sigma^2$.

Theorem 1.23 (Hoeffding's inequality). *Let X_1, \dots, X_n be independent sub-Gaussian random variables, with X_i having parameter σ_i^2 . Then \bar{X} is sub-Gaussian with parameter $\bar{\sigma}^2$. In particular, we have*

$$\max\{\mathbb{P}(\bar{X} \geq x), \mathbb{P}(\bar{X} \leq -x)\} \leq e^{-nx^2/(2\bar{\sigma}^2)}.$$

Proof. For $t \in \mathbb{R}$, we have

$$\mathbb{E}[e^{t\bar{X}}] = \mathbb{E}[e^{(t/n)\sum_i X_i}] = \prod_{i=1}^n \mathbb{E}[e^{(t/n)X_i}] \leq \prod_{i=1}^n e^{t^2\sigma_i^2/(2n^2)} = e^{t^2\bar{\sigma}^2/(2n)},$$

which shows \bar{X} is sub-Gaussian with parameter $\bar{\sigma}^2/n$. Applying part (a) of the previous proposition shows the second result. □

Remark. A direct consequence of Hoeffding's inequality is that

$$\mathbb{P}(|\bar{X}| \geq x) \leq 2e^{-nx^2/(2\bar{\sigma}^2)}.$$

The inequality is often stated in this weaker way.

Lemma 1.24 (Hoeffding's lemma). *Let X be a random variable with $\mathbb{E}X = 0$ that satisfies $a \leq X \leq b$. Then X is sub-Gaussian with parameter $(b-a)^2/4$.*

Proof. See Example Sheet 1, question 2. □

Corollary 1.25. *Let X_1, \dots, X_n be independent random variables where $\mathbb{E}[X_i] = \mu_i$ and $a_i \leq X_i \leq b_i$. Then we have*

$$\mathbb{P}(\bar{X} - \bar{\mu} \geq x) \leq \exp\left(-\frac{2n^2x^2}{\sum_{i=1}^n (b_i - a_i)^2}\right).$$

Proof. By Hoeffding's lemma, $X_i - \mu_i$ is sub-Gaussian with parameter $(b_i - a_i)^2/4$ for each i . The result now follows from theorem 1.23. □

Note that when X takes values in $[a, b]$, its variance is at most $(b - a)^2$. However, when $\text{Var}(X_i) \ll (b_i - a_i)^2$, Hoeffding's inequality can be loose (for example, when $X_i \sim \text{Bern}(p_i)$ with p_i small). In such circumstances, Bennett's or Bernstein's inequality may give better results.

Theorem 1.26 (Bennett's inequality). *Let X_1, \dots, X_n be independent random variables with $\mathbb{E}[X_i] = 0$, $\sigma_i^2 := \text{Var}(X_i) < \infty$, and $X_i \leq b$ for some $b > 0$. Define $S := \sum_{i=1}^n X_i$, $\nu := \overline{\text{Var}(X_i)}$ and $\varphi: \mathbb{R} \rightarrow \mathbb{R}$ by $\varphi(u) := e^u - 1 - u = \sum_{k=2}^{\infty} \frac{u^k}{k!}$, then for every $t > 0$ we have*

$$\log \mathbb{E}[e^{tS}] \leq \frac{n\nu}{b^2} \varphi(bt).$$

Defining $h: (0, \infty) \rightarrow [0, \infty)$ by $h(u) := (1 + u) \log(1 + u) - u$, we have for every $x > 0$ that

$$\mathbb{P}(\bar{X} \geq x) \leq \exp \left(-\frac{n\nu}{b^2} h\left(\frac{bx}{\nu}\right) \right).$$

Proof. Define $g: \mathbb{R} \rightarrow \mathbb{R}$ by

$$g(u) := \sum_{k=0}^{\infty} \frac{u^k}{(k+2)!} = \begin{cases} \frac{\varphi(u)}{u^2} & \text{if } u \neq 0, \\ \frac{1}{2} & \text{if } u = 0. \end{cases}$$

Then one can check that g is increasing on \mathbb{R} , so

$$e^{tX_i} - 1 - tX_i = t^2 X_i^2 g(tX_i) \leq t^2 X_i^2 g(tb) = X_i^2 \frac{\varphi(bt)}{b^2},$$

and therefore

$$e^{tX_i} \leq 1 + tX_i + X_i^2 \frac{\varphi(bt)}{b^2} \implies \mathbb{E}[e^{tX_i}] \leq 1 + \mathbb{E}[X_i^2] \frac{\varphi(bt)}{b^2} = 1 + \text{Var}(X_i) \frac{\varphi(bt)}{b^2}.$$

Hence for $t > 0$ we have

$$\begin{aligned} \log \mathbb{E}[e^{tS}] &= \sum_{i=1}^n \log \mathbb{E}[e^{tX_i}] \leq n \cdot \frac{1}{n} \sum_{i=1}^n \log \left(1 + \text{Var}(X_i) \frac{\varphi(bt)}{b^2} \right) \\ &\stackrel{*}{\leq} n \log \left(1 + \frac{\nu \varphi(bt)}{b^2} \right) \stackrel{**}{\leq} \frac{n\nu}{b^2} \varphi(bt). \end{aligned}$$

Here, (*) follows from the fact that \log is a concave function while (**) follows from the fact that $\log(1 + u) \leq u$ for all $u \geq 0$. This concludes the proof for the first part of the theorem.

Now, we apply the method of Chernoff bounding and find

$$\mathbb{P}(\bar{X} \geq x) = \mathbb{P}(S \geq nx) \leq \inf_{t>0} e^{-ntx} \mathbb{E}[e^{tS}] \leq \inf_{t>0} e^{-ntx + n\nu \varphi(bt)/b^2} = \exp \left(-\frac{n\nu}{b^2} h\left(\frac{bx}{\nu}\right) \right),$$

since one can check that the infimum is attained at $t = b^{-1} \log(1 + bx/\nu)$. \square

Definition 1.27. A random variable X with $\mathbb{E}X = 0$ is called *sub-Gamma in the right tail* with variance factor $\sigma^2 > 0$ and scale $c > 0$ if

$$\mathbb{E}[e^{tX}] \leq \exp \left(\frac{\sigma^2 t^2}{2(1 - ct)} \right)$$

for all $t \in [0, 1/c)$.

Note that this definition looks like that of sub-Gaussianity, except that e^{tX} can explode as t approaches $1/c$. We give some characteristics of sub-Gamma distributions:

Definition 1.28. For any $x \in \mathbb{R}$ we define $x_+ := \max(x, 0)$.

Proposition 1.29. (a) Let X be sub-Gamma in the right tail with variance factor σ^2 and scale c . Then

$$\mathbb{P}(X \geq x) \leq \exp\left(-\frac{x^2}{2(\sigma^2 + cx)}\right)$$

for all $x \geq 0$.

(b) Let X be a random variable with $\mathbb{E}X = 0$, $\mathbb{E}[X^2] \leq \sigma^2$ and $\mathbb{E}[(X_+)^q] \leq q!\sigma^2 c^{q-2}/2$ for all $q \geq 3$. Then X is sub-Gamma in the right tail with variance factor σ^2 and scale parameter c .

Proof. (a) Again, we apply a Chernoff bound: we have

$$\begin{aligned} \mathbb{P}(X \geq x) &\leq \inf_{t \in [0, 1/c)} e^{-tx} \mathbb{E}[e^{tX}] \leq \inf_{t \in [0, 1/c)} \exp\left(-tx + \frac{\sigma^2 t^2}{2(1-ct)}\right) \\ &\leq \exp\left(-\frac{x^2}{2(\sigma^2 + cx)}\right), \end{aligned}$$

where we have set $t = x/(\sigma^2 + cx) \in [0, 1/c)$ in the final step.

(b) Recall from the proof of Bennett's inequality that g is increasing and therefore for $u \leq 0$ we have $\varphi(u) = u^2 g(u) \leq u^2 g(0) = \frac{u^2}{2}$. Therefore, for every $u \in \mathbb{R}$ we have

$$\varphi(u) \leq \frac{u^2}{2} + \sum_{q=3}^{\infty} \frac{(u_+)^q}{q!}.$$

We deduce that for $t \in [0, 1/c)$ we have (note $\log(x) \leq x - 1$ for all x):

$$\log \mathbb{E}[e^{tX}] \leq \mathbb{E}(e^{tX}) - 1 = \mathbb{E}[\varphi(tX)] \leq \mathbb{E}\left[\frac{t^2 X^2}{2} + \sum_{q=3}^{\infty} \frac{t^q X_+^q}{q!}\right].$$

By Fubini's theorem, since the infinite sum has only positive terms we may interchange sum and expectation to obtain

$$\mathbb{E}\left[\frac{t^2 X^2}{2} + \sum_{q=3}^{\infty} \frac{t^q \mathbb{E}[X_+^q]}{q!}\right] = \frac{t^2 \text{Var}[X]}{2} + \sum_{q=3}^{\infty} \frac{t^q \mathbb{E}[X_+^q]}{q!} \leq \frac{\sigma^2 t^2}{2} \sum_{q=2}^{\infty} t^{q-2} c^{q-2} = \frac{\sigma^2 t^2}{2(1-ct)}.$$

□

Following this proposition, we can prove Bernstein's inequality:

Theorem 1.30 (Bernstein's inequality). Let X_1, \dots, X_n be independent random variables with $\mathbb{E}[X] = 0$, $\frac{1}{n} \sum_{i=1}^n \text{Var}(X_i) \leq \sigma^2$ and $\frac{1}{n} \sum_{i=1}^n \mathbb{E}[(X_i)_+^q] \leq q!\sigma^2 c^{q-2}/2$ some $\sigma, c > 0$ and for all $q \geq 3$. Then $S := \sum_{i=1}^n X_i$ is sub-Gamma in the right tail with variance factor $n\sigma^2$ and scale parameter c . In particular we have

$$\mathbb{P}(\bar{X} \geq x) \leq \exp\left(-\frac{nx^2}{2(\sigma^2 + cx)}\right),$$

for all $x \geq 0$.

Proof. We have by part (b) of the previous proposition

$$\log \mathbb{E}[e^{tS}] = \sum_{i=1}^n \log \mathbb{E}[e^{tX_i}] \leq n \frac{\sigma^2 t^2}{2(1-ct)},$$

and the second claim follows from part (a) of the previous proposition:

$$\mathbb{P}(\bar{X} \geq x) = \mathbb{P}(S \geq nx) \leq \exp\left(-\frac{nx^2}{2(\sigma^2 + cx)}\right).$$

□

2 Kernel density estimation

Let $X_1, \dots, X_n \stackrel{\text{iid}}{\sim} f$, and suppose we wish to estimate the density function f . The oldest way to do this is with a histogram: we divide \mathbb{R} into equally sized intervals or *bins*, and let I_x denote the bin containing $x \in \mathbb{R}$.

Definition 2.1. The *histogram density estimator* \hat{f}_n^H with bin width $b > 0$ is given by

$$\hat{f}_n^H(x) := \frac{1}{nb} \sum_{i=1}^n \mathbb{1}_{X_i \in I_x}.$$

There are a few major drawbacks to using histograms: it is difficult to choose b and the positioning of bin edges, the theoretical performance is suboptimal (mostly due to their discontinuity) and graphical display in the multivariate case is difficult.

2.1 The univariate kernel density estimator

Definition 2.2. A Borel measurable function $K: \mathbb{R} \rightarrow \mathbb{R}$ is called a *kernel* if it satisfies $\int_{\mathbb{R}} K(x) dx = 1$. A *univariate kernel density estimator* of f with kernel K and *bandwidth* $h > 0$ is defined as

$$\hat{f}_n(x) := \frac{1}{nh} \sum_{i=1}^n K\left(\frac{x - X_i}{h}\right).$$

Defining $K_h(x) := \frac{1}{h} K\left(\frac{x}{h}\right)$, we can rewrite this as

$$\hat{f}_n(x) := \frac{1}{n} \sum_{i=1}^n K_h(x - X_i).$$

Usually K is chosen to be non-negative (which ensures that K itself and \hat{f}_n are themselves density functions), and K is often chosen to be symmetric about 0. Generally, the choice of kernel K is much less important than the choice of bandwidth h .

If we consider $\hat{f}_n(x)$ as a point estimator of $f(x)$, we typically wish to minimise the *mean squared error*

$$\text{MSE}(\hat{f}_n(x)) := \mathbb{E}\left[(\hat{f}_n(x) - f(x))^2\right].$$

Other possibilities include the mean absolute error which (unlike the MSE) is scale-invariant. However, the MSE has an appealing decomposition into variance and bias terms:

$$\text{MSE}(\hat{f}_n(x)) = \text{Var}(\hat{f}_n(x)) + \text{Bias}^2(\hat{f}_n(x)).$$

We can express the MSE in terms of convolutions:

Definition 2.3. Let $g_1, g_2: \mathbb{R} \rightarrow \mathbb{R}$ be measurable. Then the *convolution* of g_1 and g_2 , denoted $g_1 * g_2$, is defined by

$$(g_1 * g_2)(x) := \int_{\mathbb{R}} g_1(x - z)g_2(z) dz.$$

We can compute

$$\begin{aligned} \text{Bias } \hat{f}_n(x) &= \mathbb{E}[\hat{f}_n(x)] - f(x) = \mathbb{E}[K_h(x - X_1)] - f(x) = \int_{\mathbb{R}} K_h(x - z)f(z) dz \\ &= (K_h * f)(x) - f(x). \end{aligned} \tag{2}$$

Analogously, letting $\xi_i := K_h(x - X_i)$ (note that these are i.i.d. random variables), we have

$$\text{Var } \hat{f}_n(x) = \frac{1}{n} \text{Var}(\xi_1) = \frac{1}{n} (\mathbb{E}[\xi_1^2] - \mathbb{E}^2[\xi_1]) = \frac{1}{n} [(K_h^2 * f)(x) - (K_h * f)^2(x)]. \quad (3)$$

To assess performance of h and K , we want to assess the performance of \hat{f}_n as an estimation of f as a function. This gives the following definition:

Definition 2.4. We define the *mean integrated squared error* or MISE as

$$\text{MISE}(\hat{f}_n) := \mathbb{E} \left(\int_{\mathbb{R}} (\hat{f}_n(x) - f(x))^2 dx \right) \stackrel{*}{=} \int_{\mathbb{R}} \text{MSE}(\hat{f}_n(x)) dx,$$

where \star follows from Fubini's theorem since the integrand is nonnegative.

We now aim to find bounds on the bias and the variance of \hat{f}_n in order to choose h and K appropriately.

2.2 Bounds on variance and bias

Definition 2.5. For a kernel K , define $R(K) := \int_{\mathbb{R}} K^2(u) du$.

Proposition 2.6. Let \hat{f}_n be the kernel density estimator with kernel K and bandwidth $h > 0$ constructed from $X_1, \dots, X_n \stackrel{\text{iid}}{\sim} f$. Then for any $x \in \mathbb{R}, h > 0, n \in \mathbb{N}$ we have

$$\text{Var } \hat{f}_n(x) \leq \frac{1}{nh} \|f\|_{\infty} R(K).$$

Proof. By eq. (3) we have

$$\begin{aligned} \text{Var } \hat{f}_n(x) &\leq \frac{1}{n} (K_h^2 * f)(x) = \frac{1}{nh^2} \int_{\mathbb{R}} K^2\left(\frac{x-z}{h}\right) f(z) dz = \frac{1}{nh} \int_{\mathbb{R}} K^2(u) f(x-uh) du \\ &\leq \frac{1}{nh} \|f\|_{\infty} \int_{\mathbb{R}} K^2(u) du = \frac{1}{nh} \|f\|_{\infty} R(K). \end{aligned} \quad (4)$$

□

Obtaining a bound on the bias is not at all straightforward: we will need to introduce conditions on both the density f and the kernel K .

Definition 2.7. Let $I \subseteq \mathbb{R}$ be an interval, fix $\beta, L > 0$, and let $m := \lceil \beta \rceil - 1$. A function $f: \mathbb{R} \rightarrow \mathbb{R}$ is said to belong to the *Hölder class* $\mathcal{H}(\beta, L)$ if f is m times differentiable on I and

$$\left| f^{(m)}(x) - f^{(m)}(y) \right| \leq L |x - y|^{\beta - m} \quad \text{for all } x, y \in I.$$

If I is unspecified, we let $I = \mathbb{R}$.

The densities in $\mathcal{H}(\beta, L)$ are denoted by

$$\mathcal{F}(\beta, L) := \left\{ f \in \mathcal{H}(\beta, L) \mid f \geq 0 \text{ and } \int_{\mathbb{R}} f dx = 1 \right\}.$$

Definition 2.8. Fix $\ell \in \mathbb{N}$. We say a kernel K is of *order* ℓ if $\int_{\mathbb{R}} x^j k(x) dx = 0$ for $j = 1, \dots, \ell - 1$.

Remark. Most kernels used in practice are of order 2. Note that a kernel of order ≥ 3 cannot be nonnegative, since we have $\int_{\mathbb{R}} x^2 K(x) dx = 0$. Therefore, the kernels are not themselves densities and the corresponding kernel density estimate is not guaranteed to be a density.

Proposition 2.9. Assume that $f \in \mathcal{F}(\beta, L)$ and that K is a kernel of order $\ell := \lceil \beta \rceil$, and furthermore assume that

$$\mu_\beta(K) := \int_{\mathbb{R}} |u|^\beta |K(u)| \, du < \infty.$$

Then the kernel density estimate with bandwidth h and kernel K based on $X_1, \dots, X_n \sim f$ satisfies

$$\left| \text{Bias } \hat{f}_n(x) \right| \leq \frac{L}{(\ell-1)!} \mu_\beta(K) h^\beta \quad \text{for all } x \in \mathbb{R}, h > 0, n \in \mathbb{N}.$$

Proof. By eq. (2), we have

$$\text{Bias } \hat{f}_n(x) = \frac{1}{h} \int_{\mathbb{R}} K\left(\frac{x-z}{h}\right) f(z) \, dz - f(x) = \int_{\mathbb{R}} K(u) (f(x-uh) - f(x)) \, du.$$

By applying Taylor's theorem with the Lagrange remainder we obtain, with $m = \lceil \beta \rceil - 1$, that

$$f(x-uh) - f(x) = \sum_{j=1}^{m-1} \frac{(-uh)^j}{j!} f^{(j)}(x) + \frac{(-uh)^m}{m!} f^{(m)}(x - \tau uh) \quad \text{for some } \tau \in [0, 1].$$

Since $\int_{\mathbb{R}} u^j K(u) \, du = 0$ for all $j \leq m$, plugging the sum into the integral will give 0. Therefore, we find

$$\text{Bias } \hat{f}_n(x) = \frac{(-h)^m}{m!} \int_{\mathbb{R}} u^m K(u) f^{(m)}(x - \tau uh) \, du = \frac{(-h)^m}{m!} \int_{\mathbb{R}} u^m K(u) \left[f^{(m)}(x - \tau uh) - f^{(m)}(x) \right] \, du,$$

where the last inequality follows again from the fact that K is of order $m+1$.

Now we use that $f \in \mathcal{F}(\beta, L)$, and conclude

$$\left| \text{Bias } \hat{f}_n(x) \right| \leq \frac{Lh^m}{m!} \int_{\mathbb{R}} |u|^m |K(u)| |\tau uh|^{\beta-m} \, du \leq \frac{Lh^\beta}{m!} \int_{\mathbb{R}} |u|^\beta |K(u)| \, du = \frac{L}{(\ell-1)!} \mu_\beta(K) h^\beta,$$

which concludes the proof. \square

Combining propositions 2.6 and 2.9, we find that

$$\text{MSE } \hat{f}_n(x) \leq \frac{1}{nh} R(K) \|f\|_\infty + \frac{L^2}{((\ell-1)!)^2} \mu_\beta^2(K) h^{2\beta}.$$

By minimising this w.r.t. h , we find that the optimal h is given by

$$h_n^* = \left(\frac{((\ell-1)!)^2 \|f\|_\infty R(K)}{2\beta L^2 \mu_\beta^2(K)} \right)^{1/(2\beta+1)} n^{-1/(2\beta+1)},$$

and the corresponding MSE is given by

$$\text{MSE } \hat{f}_n(x) \leq \left(\frac{\|f\|_\infty^{2\beta} R(K)^{2\beta} L^2 \mu_\beta^2(K) [(2\beta)^{2\beta+1} + 1]}{((\ell-1)!)^2 (2\beta)^{2\beta}} \right) n^{-2\beta/(2\beta+1)},$$

This $O(n^{-2\beta/(2\beta+1)})$ bound on the rate is slower than the $O(1/n)$ rate found in parametric problems, but such a rate is only obtained when the assumed model is correct.

We can strengthen this as follows:

Theorem 2.10. Assume that K is a kernel of order $\ell := \lceil \beta \rceil$ and that $\mu_\beta(K)$ and $R(K)$ are both finite. Fix $\alpha > 0$ and let $h = \alpha n^{-1/(2\beta+1)}$. Then there exists $C > 0$, independent of n , such that

$$\sup_{x \in \mathbb{R}} \sup_{f \in \mathcal{F}(\beta, L)} \text{MSE } \hat{f}_n(x) \leq C n^{-2\beta/(2\beta+1)}.$$

Proof. We will show that the class $\mathcal{F}(\beta, L)$ is uniformly bounded in supremum norm. Let K^* be a bounded kernel of order ℓ (see example sheet **TODO:**), then by the previous proposition with $h = 1$ we have by nonnegativity of f that

$$\begin{aligned} f(x) &\leq \left| f(x) - \int_{-\infty}^{\infty} K^*(x-z)f(z) dz \right| + \left| \int_{-\infty}^{\infty} K^*(x-z)f(z) dz \right| \\ &\leq \left| \text{Bias } \hat{f}_{n,K^*}(x) \right| + \|K^*\|_{\infty} \int_{-\infty}^{\infty} f(z) dz \\ &\leq \frac{L}{(\ell-1)!} \mu_{\beta}(K^*) + \|K^*\|_{\infty}, \end{aligned}$$

and this bound is independent of f and x .

Now we have

$$\text{MSE } \hat{f}_n(x) \leq \frac{R(K)\|f\|_{\infty}}{nh} + \frac{L^2}{((\ell-1)!)^2} \mu_{\beta}^2(K) h^{2\beta} \leq C n^{-2\beta/(2\beta+1)}.$$

□

2.3 Bounds on the integrated variance and bias

To bound the MISE, we will give bounds on the integrated variance and bias.

Proposition 2.11. *Let \hat{f}_n denote the kernel density estimate with bandwidth h and kernel K based on $X_1, \dots, X_n \stackrel{\text{iid}}{\sim} P$ (where P is a distribution on \mathbb{R}). Then*

$$\int_{-\infty}^{\infty} \text{Var } \hat{f}_n(x) dx = \frac{1}{nh} R(K).$$

Proof. We have by Fubini and eq. (4) that

$$\begin{aligned} \int_{-\infty}^{\infty} \text{Var } \hat{f}_n(x) dx &\leq \frac{1}{nh^2} \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} K^2\left(\frac{x-z}{h}\right) f(z) dz dx = \frac{1}{nh^2} \int_{-\infty}^{\infty} f(z) \int_{-\infty}^{\infty} K^2\left(\frac{x-z}{h}\right) dx dz \\ &= \frac{1}{nh} R(K) \int_{-\infty}^{\infty} f(z) dz = \frac{1}{nh} R(K). \end{aligned}$$

□

Recap 2.12. Let $[a, b] = I \subseteq \mathbb{R}$ be an interval, then $f: I \rightarrow \mathbb{R}$ is called *absolutely continuous* if, for every $\varepsilon > 0$, there exists $\delta > 0$ such that, whenever $(x_1, y_1), \dots, (x_m, y_m)$ are disjoint subintervals of I with $\sum_{i=1}^m (y_i - x_i) < \delta$, we have $\sum_{i=1}^m |f(y_i) - f(x_i)| < \varepsilon$.

It is known that absolute continuity is equivalent to being differentiable Lebesgue almost everywhere with a so-called *weak derivative* f' that satisfies $f(x) = f(a) + \int_a^x f'(t) dt$ for all $x \in [a, b]$.

Recap 2.13. The *generalised Minkowski inequality* states that any Borel measurable function $g: \mathbb{R}^2 \rightarrow \mathbb{R}$ we have that

$$\int_{\mathbb{R}} \left(\int_{\mathbb{R}} g(u, x) du \right)^2 dx \leq \left(\int_{\mathbb{R}} \left(\int_{\mathbb{R}} g^2(u, x) dx \right)^{1/2} du \right)^2.$$

To obtain bounds on the integrated squared bias, we will require smoothness conditions w.r.t. the $L^2(\mathbb{R})$ norm.

Definition 2.14. Fix $\beta, L > 0$ and let $m := \lceil \beta \rceil - 1$. The *Nikolski class* $\mathcal{N}(\beta, L)$ consists of functions $f: \mathbb{R} \rightarrow \mathbb{R}$ that are $(m-1)$ times differentiable and for which $f^{(m-1)}$ is absolutely continuous with weak derivative $f^{(m)}$ satisfying

$$\left\{ \int_{-\infty}^{\infty} \left[f^{(m)}(x+t) - f^{(m)}(x) \right]^2 dx \right\}^{1/2} \leq L|t|^{\beta-m} \quad \text{for all } t \in \mathbb{R}.$$

The densities in $\mathcal{N}(\beta, L)$ are denoted by $\mathcal{F}_{\mathcal{N}}(\beta, L)$.

Proposition 2.15. Fix $\beta, L > 0$ and let K be a kernel of order $\ell := \lceil \beta \rceil$. Let \hat{f}_n denote the KDE with kernel K and bandwidth h based on $X_1, \dots, X_n \stackrel{\text{iid}}{\sim} f \in \mathcal{F}_{\mathcal{N}}(\beta, L)$. Then we have

$$\int_{-\infty}^{\infty} \text{Bias}^2 \hat{f}_n(x) dx \leq \frac{L^2}{((\ell-1)!)^2} \mu_{\beta}^2(K) h^{2\beta}.$$

Proof. **TODO:** write this out (integration + taylor expansion + 2x minkowski). □

Putting everything together, we obtain the following:

Theorem 2.16. Fix $\beta, L > 0$, and let K be a kernel of order $\ell = \lceil \beta \rceil$ with $R(K)$ and $\mu_{\beta}(K)$ finite. Let \hat{f}_n be the KDE with kernel K and bandwidth h based on $X_1, \dots, X_n \stackrel{\text{iid}}{\sim} f \in \mathcal{F}_{\mathcal{N}}(\beta, L)$. Then we have

$$\text{MISE } \hat{f}_n \leq \frac{R(K)}{nh} + \frac{L^2}{((\ell-1)!)^2} \mu_{\beta}^2(K) h^{2\beta}.$$

In particular, fixing $\alpha > 0$ and taking $h = \alpha n^{-1/(2\beta+1)}$, there exists $C > 0$ independent of n such that

$$\sup_{f \in \mathcal{F}_{\mathcal{N}}(\beta, L)} \text{MISE } \hat{f}_{n,h,K} \leq C n^{-2\beta/(2\beta+1)}.$$

2.4 Bandwidth selection

The choice of bandwidth in the previous theorem is not practical since we have not specified α and β is typically unknown.

2.4.1 Least squares cross validation

One possible approach is *least squares cross validation*. For this, note that minimising the MISE is equivalent to minimising

$$\text{MISE}(\hat{f}_n) - \int_{\mathbb{R}} f^2(x) dx = \mathbb{E} \left[\int_{\mathbb{R}} \hat{f}_n^2(x) dx \right] - 2\mathbb{E} \left[\int_{\mathbb{R}} \hat{f}_n(x) f(x) dx \right],$$

and it can be checked that an unbiased estimator for the above is given by

$$\text{LSCV}(h) := \int_{\mathbb{R}} \hat{f}_n^2(x) dx - \frac{2}{n} \sum_{i=1}^n \hat{f}_{n,-i}(X_i),$$

where $\hat{f}_{n,-i}$ is the KDE based on all observations except X_i . We now choose h such that LSCV is minimised.

2.4.2 Lepski

TODO: write this subsection (**TODO:** understand this first)

2.5 Choice of kernel

To choose a kernel, we first fix the scale of the kernel by setting $\mu_2(K) = 1$. Now, by our bound on the MISE (theorem 2.16) it is reasonable to minimise $R(K)$, where for simplicity we assume that K is a nonnegative second-order kernel. The solution is the Epanechnikov kernel

$$K_E(x) := \frac{3}{4\sqrt{5}} \left(1 - \frac{x^2}{5}\right) \mathbb{1}_{|x| \leq \sqrt{5}},$$

and the ratio $R(K_E)/R(K)$ is called the *efficiency* of a kernel K . We find that for different kernels, the efficiency is greater than 0.9, which shows that kernel shape does not affect efficiency greatly.

2.6 Multivariate density estimation

The general d -dimensional KDE is

$$\hat{f}_n(x) := \frac{1}{n\sqrt{\det H}} \sum_{i=1}^n K(H^{-1/2}(x - X_i)),$$

where H is a symmetric positive-definite bandwidth matrix (often chosen to be diagonal or a multiple of I). If $H = h^2 I$, it can be shown that, under an appropriate definition of a β smoothness class, we have an optimal bandwidth of order $n^{-1/(d+2\beta)}$, and with this choice, a MISE of order $n^{-2\beta/(d+2\beta)}$. This is called the “curse of dimensionality”: the higher the dimension becomes, the harder nonparametric estimation gets.

3 Nonparametric regression

3.1 Fixed and random design

In *fixed design*, we assume we have data $x_1 \leq \dots \leq x_n$ and the response variables satisfy

$$Y_i := m(x_i) + \sqrt{v(x_i)}\varepsilon_i,$$

where the ε_i are independent, mean-zero random variables with $\text{Var}(\varepsilon_i) = 1$. The function m is called the *regression function*, and the function v is the *variance function*. If v is constant, the model is called *homoscedastic*, else it is called *heteroscedastic*.

In *random design*, we assume we have i.i.d. data pairs (X_i, Y_i) with

$$Y_i = m(X_i) + \sqrt{v(X_i)}\varepsilon_i,$$

where the ε_i are again independent with $\mathbb{E}[\varepsilon_1|X_1] = 0$ and $\text{Var}(\varepsilon_1|X_1) = 1$. The regression function is given by $m(x) = \mathbb{E}(Y_1|X_1 = x)$ and the variance function by $v(x) = \text{Var}(Y_1|X_1 = x)$.

3.2 Local polynomial estimators

We will assume the fixed design setting.

Definition 3.1. Let K be a kernel, $h > 0$ a bandwidth and $p \in \mathbb{N}$. Then the *local polynomial estimator of degree p with bandwidth h and kernel K* , denoted $\hat{m}_n(\cdot; p) \equiv \hat{m}_n(\cdot; p, h, K)$, is constructed at $x \in \mathbb{R}$ by fitting a polynomial p to the data using weighted least squares, where the pair (x_i, Y_i) is assigned weight $K_h(x_i - x)$.

To write this in formulas, define $Q(u) := (1, u, \frac{u^2}{2}, \dots, \frac{u^p}{p!}) \in \mathbb{R}^{p+1}$ and $Q_h(\cdot) = Q(\cdot/h)$, we then have

$$\hat{m}_n(x; p) = \hat{\beta}_0, \quad \text{where } \hat{\beta} := \arg \min_{\beta \in \mathbb{R}^{p+1}} \sum_{i=1}^n (Y_i - \beta^\top Q_h(x_i - x))^2 K_h(x_i - x).$$

In matrix-vector notation, writing

$$X \equiv X(x; p, h) := \begin{pmatrix} Q_h(x_1 - x)^\top \\ \vdots \\ Q_h(x_n - x)^\top \end{pmatrix} \in \mathbb{R}^{n \times (p+1)}, \quad Y := \begin{pmatrix} Y_1 \\ \vdots \\ Y_n \end{pmatrix} \in \mathbb{R}^n,$$

$$W \equiv W(x; h, K) := \text{diag}(K_h(x_1 - x), \dots, K_h(x_n - x)) \in \mathbb{R}^{n \times n},$$

we have

$$\hat{\beta} = \arg \min_{\beta \in \mathbb{R}^{p+1}} (Y - X\beta)^\top W (Y - X\beta).$$

By standard weighted least squares theory, we know that $\hat{\beta}$ must satisfy $X^\top W X \hat{\beta} = X^\top W Y$.

Proposition 3.2. Suppose $X^\top W X$ is positive definite. Then we have

$$\hat{\beta} = (X^\top W X)^{-1} X^\top W Y.$$

We will assume from here on that $X^\top W X$ is indeed positive definite. In this case, since the entries of W and X are functions of $x_i - x$, we can write the local polynomial estimator in the form

$$\hat{m}_n(x) = n^{-1} \sum_{i=1}^n w(x, x_i) Y_i.$$

(?) The set of weights $\{w(x, x_i)\}$ is called the *effective kernel* at x .

For $p = 0$ and $p = 1$, there exist explicit formulas for the local polynomial estimator of degree p .

Proposition 3.3 (Reproducing property). *Let $\{w_p(x, x_i)\}$ denote the effective kernel of a local polynomial estimator of degree p based on data $(x_1, Y_1), \dots, (x_n, Y_n)$, and let R denote a polynomial of degree at most p . If $X^\top W X$ is positive definite, then*

$$\frac{1}{n} \sum_{i=1}^n w_p(x, x_i) R(x_i) = R(x).$$

Proof. See example sheet 2 question 3. □

Before we can study the bias and variance of local polynomial estimators, we require the following lemma:

Lemma 3.4. *Let K be a kernel that vanishes outside $[-1, 1]$, and suppose that $n^{-1} X^\top W X$ is positive definite with minimal eigenvalue $\lambda_0 \equiv \lambda_{0,n,x} > 0$. Then*

- (i) $\sup_{x \in [0,1]} \max_{i=1, \dots, n} \frac{1}{n} |w(x, x_i)| \leq \frac{2\|K\|_\infty}{\lambda_0 n h};$
- (ii) $n^{-1} \sum_{i=1}^n |w(x, x_i)| \leq \frac{2\|K\|_\infty}{\lambda_0 n h} \cdot \sum_{i=1}^n \mathbb{1}_{|x_i - x| \leq h};$
- (iii) $w(x, x_i) = 0$ when $|x_i - x| > h$.

Proof. (i) Note that $n^{-1} w(x, x_i)$ is the $(0, i)$ entry of the matrix $(X^\top W X)^{-1} X^\top W$, and it is therefore less than the norm of the i -th column of $(X^\top W X)^{-1} X^\top W$. For $x \in [0, 1]$ and $i = 1, \dots, n$, we find

$$\begin{aligned} \frac{1}{n} |w(x, x_i)| &\leq \|(X^\top W X)^{-1} Q_h(x_i - x) K_h(x_i - x)\| \stackrel{*}{\leq} \|K_h\|_\infty \|(X^\top W X)^{-1}\| \|Q_h(x_i - x)\| \mathbb{1}_{|x_i - x| \leq h} \\ &= \frac{\|K\|_\infty}{h} \frac{1}{\lambda_0 n} \|Q_h(x_i - x)\| \mathbb{1}_{|x_i - x| \leq h} \leq \frac{\|K\|_\infty}{\lambda_0 n h} \|Q(1)\| = \frac{\|K\|_\infty}{\lambda_0 n h} \left(\sum_{j=0}^p \frac{1}{(j!)^2} \right)^{1/2} \\ &\leq \frac{\|K\|_\infty}{\lambda_0 n h} e^{1/2} \leq \frac{2\|K\|_\infty}{\lambda_0 n h}. \end{aligned}$$

Here, the indicator function in \star appears because K vanishes outside $[-1, 1]$. Since the upper bound is independent of both x and i , the claim is proven.

(ii) Similarly as before, we find

$$\frac{1}{n} \sum_{i=1}^n |w(x, x_i)| \leq \frac{\|K\|_\infty}{\lambda_0 n h} \sum_{i=1}^n \|Q_h(x_i - x)\| \mathbb{1}_{|x_i - x| \leq h} \leq \frac{2\|K\|_\infty}{\lambda_0 n h} \sum_{i=1}^n \mathbb{1}_{|x_i - x| \leq h}.$$

(iii) This follows immediately from inequality \star in the proof of (i). □

Now, we can compute bounds on the variance and bias of our local polynomial estimator. For simplicity, we will assume $x_i = i/n$.

Proposition 3.5. *Assume the model $Y_i = m(x_i) + v^{1/2} \varepsilon_i$ with $m \in \mathcal{H}(\beta, L)$ on $[0, 1]$ and $\max_i v(x_i) \leq \sigma_{\max}^2$. Let K be a kernel that vanishes outside $[-1, 1]$, and suppose that λ_0 , the minimal eigenvalue of $n^{-1} X^\top W X$, is strictly positive. Then, for $p \geq \lceil \beta \rceil - 1 =: \beta_0$ and for each $x_0 \in [0, 1]$, $n \in \mathbb{N}$ and $h \geq 1/(2n)$, we have*

$$\text{Var } \hat{m}_n(x_0; p) \leq \frac{16\|K\|_\infty^2 \sigma_{\max}^2}{\lambda_0^2 n h}, \quad |\text{Bias } \hat{m}_n(x_0; p)| \leq \frac{8L\|K\|_\infty}{\lambda_0 \beta_0!} h^\beta.$$

Proof. Using the previous lemma, we obtain

$$\begin{aligned}\text{Var } \hat{m}_n(x_0; p) &= \text{Var} \left(\frac{1}{n} \sum_{i=1}^n w(x_0, x_i) Y_i \right) = \frac{1}{n^2} \sum_{i=1}^n w(x_0, x_i)^2 \text{Var}(Y_i) \\ &\leq \frac{\sigma_{\max}^2}{n^2} \sum_{i=1}^n w(x_0, x_i)^2 \leq \sigma_{\max}^2 \left(\sup_{x \in [0,1]} \max_i \frac{1}{n} |w(x, x_i)| \right) \frac{1}{n} \sum_{i=1}^n |w(x_0, x_i)| \\ &\leq \frac{4\|K\|_{\infty}^2 \sigma_{\max}^2}{\lambda_0^2 n^2 h^2} \sum_{i=1}^n \mathbb{1}_{|i/n - x_0| \leq h}.\end{aligned}$$

Now, we have $|i/n - x_0| \leq h \iff i \in [nx_0 - nh, nx_0 + nh]$, and there are at most $2nh + 1$ integers in this interval. Recalling that $1 \leq 2nh$ we obtain

$$\text{Var } \hat{m}_n(x_0; p) \leq \frac{4(2nh + 1)\|K\|_{\infty} \sigma_{\max}^2}{\lambda_0^2 n^2 h^2} \leq \frac{16\|K\|_{\infty} \sigma_{\max}^2}{\lambda_0^2 nh}.$$

For the bias, we will first use a Taylor expansion combined with the reproducing property proposition 3.3. Firstly, since $\frac{1}{n} \sum_{i=1}^n w(x_0, x_i) = 1$ by this property, we have

$$\text{Bias } \hat{m}_n(x_0; p) = \left(\frac{1}{n} \sum_{i=1}^n w(x_0, x_i) m(x_i) \right) - m(x_0) = \frac{1}{n} \sum_{i=1}^n w(x_0, x_i) \{m(x_i) - m(x_0)\}.$$

Now, we apply a Taylor expansion to write

$$m(x_i) - m(x_0) = P(x_i - x_0) + \frac{1}{\beta_0!} m^{(\beta_0)}(x_0 + \tau_i(x_i - x_0))(x_i - x_0)^{\beta_0},$$

where P has degree at most $\beta_0 - 1 < p$ and a constant coefficient equal to 0, and $\tau_i \in [0, 1]$. By the reproducing property we have $n^{-1} \sum_{i=1}^n w(x_0, x_i) P(x_i - x_0) = P(0) = 0$, so we obtain

$$\begin{aligned}\text{Bias } \hat{m}_n(x; p) &= \frac{1}{n} w(x_0, x_i) \sum_{i=1}^n w(x_0, x_i) \frac{m^{(\beta_0)}(x_0 + \tau_i(x_i - x_0))}{\beta_0!} (x_i - x_0)^{\beta_0} \\ &= \frac{1}{n} w(x_0, x_i) \sum_{i=1}^n w(x_0, x_i) \frac{m^{(\beta_0)}(x_0 + \tau_i(x_i - x_0)) - m^{(\beta_0)}(x_0)}{\beta_0!} (x_i - x_0)^{\beta_0},\end{aligned}$$

where the last line follows again from the reproducing property. Now we apply the fact that $m \in \mathcal{H}(\beta, L)$ with the previous lemma and find

$$\begin{aligned}|\text{Bias } \hat{m}_n(x; p)| &\leq \frac{L}{n} \sum_{i=1}^n |w(x_0, x_i)| \frac{|x_i - x_0|^{\beta}}{\beta_0!} = \frac{L}{n} \sum_{i=1}^n |w(x_0, x_i)| \frac{|x_i - x_0|^{\beta}}{\beta_0!} \mathbb{1}_{|x_i - x_0| \leq h} \\ &\leq \frac{Lh^{\beta}}{n\beta_0!} \sum_{i=1}^n |w(x_0, x_i)| \leq \frac{8L\|K\|_{\infty}}{\lambda_0\beta_0!} h^{\beta}.\end{aligned}$$

□

So again, we have a variance term of order $1/nh$ and a bias term of order h^{β} . However, both our bounds depend on λ_0 , which depends on both n and x .

Proposition 3.6. *Suppose $x_i = i/n$ and that $K(u) \geq K_0 \mathbb{1}_{|u| \leq \Delta}$ for some $K_0, \Delta > 0$. Then, for $n \geq 2$ and $h \leq \frac{1}{4\Delta}$,*

$$\inf_{x \in [0,1]} \lambda_{0,n,x} \geq K_0 \inf_{v \in S^p} \min \left\{ \int_0^{\Delta} (v^{\top} Q(u))^2 du, \int_{-\Delta}^0 (v^{\top} Q(u))^2 du \right\} - \frac{(4\Delta + 2)K_0 e^{\Delta^2}}{nh}.$$

Proof. Since $\lambda_{0,n,x} = \inf_{v \in S^p} v^\top \left(\frac{1}{n} X^\top W X \right) v$, we will try to bound this quantity from below.

Let $u_i := \frac{x_i - x}{h}$ for $i = 1, \dots, n$, so $u_1 \leq \frac{x_1}{h} = \frac{1}{nh}$, and let $u_0 := 0$.

First, we assume $x < 1 - h\Delta$, so that $u_n > \frac{1 - (1 - h\Delta)}{h} = \Delta$. Then, for any $v \in S^p$, we have

$$\begin{aligned} v^\top \left(\frac{1}{n} X^\top W X \right) &= \frac{1}{n} (Xv)^\top W (Xv) = \frac{1}{n} \sum_{i=1}^n (Xv)_i^2 W_{ii} \\ &= \frac{1}{nh} \sum_{i=1}^n (v^\top Q(u_i))^2 K(u_i) \geq \frac{K_0}{nh} \sum_{i=1}^n (v^\top Q(u_i))^2 \mathbb{1}_{u_i \in [0, \Delta]}. \end{aligned}$$

This is a Riemann sum (since $u_{i+1} - u_i = 1/(nh)$) and we will try to approximate it by the corresponding integral $K_0 \int_0^\Delta (v^\top Q(u))^2 du$.

To do this, note that we have for $u \in [0, \Delta]$ that

$$\|Q(u)\| \leq \left\{ \sum_{\ell=0}^p \frac{u^{2\ell}}{(\ell!)} \right\}^{1/2} \leq \left(\sum_{\ell=0}^\infty \frac{u^{2\ell}}{\ell!} \right)^{1/2} \leq e^{\Delta^2/2}.$$

Furthermore, for $a, b \geq 0$, it can be shown that

$$|b^\ell - a^\ell| \leq \max(1, \ell/2) (b^{\ell-1} + a^{\ell-1}) |b - a|,$$

and using this we obtain for $u, v \in [0, \Delta]$ that

$$\begin{aligned} \|Q(u) - Q(v)\| &\leq \left\{ \sum_{\ell=1}^p \frac{\max(1, \ell^2/4) (u^{\ell-1} + v^{\ell-1})^2 (u - v)^2}{(\ell!)^2} \right\}^{1/2} \\ &\leq |u - v| \left\{ \sum_{\ell=1}^p \frac{\max(1, \ell^2/4) (2\Delta^{\ell-1})^2}{(\ell!)^2} \right\}^{1/2} \leq |u - v| \left\{ \sum_{\ell=1}^p \frac{\max(4, \ell^2) \Delta^{2\ell-2}}{(\ell!)^2} \right\}^{1/2} \\ &\stackrel{*}{\leq} 2|u - v| \left\{ \sum_{\ell=0}^{p-1} \frac{\Delta^{2\ell}}{(\ell!)^2} \right\}^{1/2} \leq 2|u - v| \left\{ \sum_{\ell=0}^\infty \frac{\Delta^{2\ell}}{\ell!} \right\}^{1/2} = 2e^{\Delta^2/2} |u - v|, \end{aligned}$$

where in \star we used that $\max(4, \ell^2) \leq 4\ell^2$.

Now we estimate the difference between the Riemann sum and the corresponding integral. Write $i_- := \min \{i : u_i > 0\}$ and $i_+ := \min \{i : u_i > \Delta\}$. Then

$$\begin{aligned} \frac{1}{nh} \sum_{i=1}^n (v^\top Q(u_i))^2 \mathbb{1}_{u_i \in [0, \Delta]} &= \sum_{i=i_-}^{i_+-1} \int_{u_{i-1}}^{u_i} (v^\top Q(u_i))^2 du \\ &\leq \sum_{i=i_-}^{i_+} \int_{\max(u_{i-1}, 0)}^{\min(u_i, \Delta)} (v^\top Q(u_i))^2 du + \frac{1}{nh} \sup_{u \in [0, \Delta]} (v^\top Q(u))^2. \end{aligned}$$

(Note that the last integral, from u_{i-1} to Δ , is 0 since) Therefore we obtain

$$\begin{aligned} &\left| \frac{1}{nh} \sum_{i=1}^n (v^\top Q(u_i))^2 \mathbb{1}_{u_i \in [0, \Delta]} - \int_0^\Delta (v^\top Q(u))^2 du \right| \\ &\leq \left| \sum_{i=i_-}^{i_+} \int_{\max(u_{i-1}, 0)}^{\min(u_i, \Delta)} (v^\top Q(u_i))^2 - (v^\top Q(u))^2 du \right| + \frac{1}{nh} \sup_{u \in [0, \Delta]} (v^\top Q(u))^2. \end{aligned} \tag{5}$$

Now note that by Cauchy-Schwarz since $\|v\| = 1$ we have

$$\begin{aligned} |(v^\top Q(u_i))^2 - (v^\top Q(u))^2| &= |(v^\top Q(u_i) + v^\top Q(u))|(v^\top Q(u_i) - v^\top Q(u))| \\ &\leq (\|Q(u_i)\| + \|Q(u)\|)\|Q(u_i) - Q(u)\| \\ &\leq 2e^{\Delta^2/2} \cdot 2e^{\Delta^2/2}|u_i - u| = 4e^{\Delta^2}|u_i - u|, \end{aligned}$$

and therefore

$$(5) \leq 4e^{\Delta^2} \sum_{i=i_-}^{i_+} \int_{\max(u_{i-1}, 0)}^{\min(u_i, \Delta)} (u_i - u) du + \frac{e^{\Delta^2}}{nh} \leq \frac{(4\Delta + 1)e^{\Delta^2}}{nh},$$

which concludes the case $x < 1 - h\Delta$.

Suppose $x \geq 1 - h\Delta$, then we have $u_1 \leq -\Delta$ and $u_n \geq 0$, and we can apply very similar arguments to reach the desired conclusion. \square

Now that we have bounds on the variance and bias, we can prove our uniform bound:

Theorem 3.7. *Under the conditions of the previous two propositions, if we choose $h = \alpha n^{-1/(2\beta+1)}$ for some $\alpha > 0$, there exists $n_0 \in \mathbb{N}, C > 0$, depending only on $\beta, L, \alpha, K, \sigma_{\max}^2$, such that*

$$\sup_{m \in \mathcal{H}(\beta, L)} \sup_{x_0 \in [0, 1]} \mathbb{E} \left[\{\hat{m}_n(x_0; p) - m(x_0)\}^2 \right] \leq C n^{-2\beta/(2\beta+1)}.$$

3.3 Splines

3.3.1 Cubic splines

Let $n \geq 3$ and $a \leq x_1 < \dots < x_n \leq b$.

Definition 3.8. A function $g: [a, b] \rightarrow \mathbb{R}$ is called a *cubic spline* with *knots* x_1, \dots, x_n if

1. g is a cubic polynomial on each interval $(a, x_1), (x_1, x_2), \dots, (x_n, b)$;
2. $g \in C^2[a, b]$.

Furthermore, g is called *natural* if it is linear on $[a, x_1]$ and $[x_n, b]$, (i.e., $g''(a) = g'''(a) = g''(b) = g'''(b) = 0$).

We often represent a natural cubic spline by the vectors $\mathbf{g} \in \mathbb{R}^n$ with $g_i = g(x_i)$, and $\boldsymbol{\gamma} \in \mathbb{R}^{n-2}$ with $\gamma_i = g''(x_i)$ (excluding γ_1 and γ_n). Writing $h_i := x_{i+1} - x_i$ we have for $x \in [x_i, x_{i+1}]$ that

$$g(x) = \frac{(x - x_i)g_{i+1} - (x_{i+1} - x)g_i}{h_i} - \frac{1}{6}(x - x_i)(x_{i+1} - x) \left\{ \left(1 + \frac{x - x_i}{h_i} \right) \gamma_{i+1} + \left(1 + \frac{x_{i+1} - x}{h_i} \right) \gamma_i \right\}.$$

Proposition 3.9. *Given $\mathbf{g} \in \mathbb{R}^n$, there exists a unique natural cubic spline g with knots at x_1, \dots, x_n satisfying $g(x_i) = g_i$ for all i , and there exists $K \succeq 0$ (depending on x_1, \dots, x_n) such that*

$$\int_a^b g''(x)^2 dx = \mathbf{g}^\top K \mathbf{g}.$$

We call g the *natural cubic spline interpolant to \mathbf{g} at x_1, \dots, x_n* .

Definition 3.10. We define $\mathcal{S}_2[a, b]$ as the set of real-valued functions on $[a, b]$ with an absolutely continuous first derivative. For $f \in \mathcal{S}_2[a, b]$, we define the *roughness* of f by $R(f) := \int_a^b f''(x)^2 dx$.

Proposition 3.11. *For any $\mathbf{g} \in \mathbb{R}^n$, the natural cubic spline interpolant to \mathbf{g} at x_1, \dots, x_n is the unique minimiser of R over all $g \in \mathcal{S}_2[a, b]$ that satisfy $g(x_i) = g_i$ for all i .*

3.3.2 Natural cubic smoothing splines

Consider the nonparametric regression model $Y_i = g(x_i) + \sigma\varepsilon_i$, where $\mathbb{E}[\varepsilon_i] = 0$ and $\text{Var}[\varepsilon_i] = 1$. A way to estimate a nonparametric regression function is to balance data fidelity against roughness of the curve, which can be done by minimising

$$S_\lambda(g) := \sum_{i=1}^n (Y_i - g(x_i))^2 + \lambda R(g),$$

where $\lambda > 0$ is a regularisation parameter. For small λ , this is almost an exact fit to the data. For large λ , we are basically minimising $|g''|$, which means we will approximate the linear regression fit.

Theorem 3.12. *For each $\lambda \in (0, \infty)$, there exists a unique minimiser \hat{g}_λ of S_λ over $\mathcal{S}_2[a, b]$. It is the natural cubic spline with knots at x_1, \dots, x_n and $\mathbf{g} = (I + \lambda K)^{-1} \mathbf{Y}$.*

Proof. If g is not a natural cubic spline, we know that the natural cubic spline g^* which interpolates $g(x_1), \dots, g(x_n)$ at x_1, \dots, x_n has a strictly lower value of S_λ , so we know the minimiser must be a natural cubic spline.

If g is a natural cubic spline, then there exists $K \succeq 0$ such that

$$\begin{aligned} S_\lambda(g) &= (\mathbf{Y} - \mathbf{g})^\top (\mathbf{Y} - \mathbf{g}) + \lambda \mathbf{g}^\top K \mathbf{g} \\ &= \mathbf{g}^\top (I + \lambda K) \mathbf{g} - 2\mathbf{Y}^\top \mathbf{g} + \mathbf{Y}^\top \mathbf{Y}, \end{aligned}$$

and by “completing the square” we write, for some Z independent of \mathbf{g} ,

$$S_\lambda(g) = (\mathbf{g} - (I + \lambda K)^{-1} \mathbf{Y})^\top (I + \lambda K) (\mathbf{g} - (I + \lambda K)^{-1} \mathbf{Y}) + Z$$

Since $I + \lambda K$ is positive definite it follows that $\mathbf{g} = (I + \lambda K)^{-1} \mathbf{Y}$ gives the minimiser. \square

The function \hat{g}_λ is called the *natural cubic smoothing spline* for data $(x_1, Y_1), \dots, (x_n, Y_n)$ and smoothing parameter λ .

3.3.3 Choice of smoothing parameter

We are left with the question of how to choose the smoothing parameter λ . A standard method is to minimise the cross-validation score

$$\text{CV}(\lambda) := \frac{1}{n} \sum_{i=1}^n (Y_i - \hat{g}_{-i, \lambda}(x_i))^2,$$

where $\hat{g}_{-i, \lambda}$ is the natural cubic smoothing spline for all data points except (x_i, Y_i) . It seems like computing $\text{CV}(\lambda)$ requires the computation of n natural cubic smoothing splines, but it turns out that this is not the case:

Proposition 3.13. *Write $A(\lambda) = (I + \lambda K)^{-1}$, then we have*

$$\text{CV}(\lambda) = \frac{1}{n} \sum_{i=1}^n \left(\frac{Y_i - \hat{g}_\lambda(x_i)}{1 - A_{ii}(\lambda)} \right)^2.$$

Proof. Example sheet 3. \square

In the above formula, we can consider the quantity $A_{ii}(\lambda)$ as the “leverage” of the i -th observation. In the *generalised cross-validation* score, we give every observation equal leverage: it is defined as

$$\text{GCV}(\lambda) := \frac{1}{n} \sum_{i=1}^n \left(\frac{Y_i - \hat{g}_\lambda(x_i)}{1 - n^{-1} \text{tr } A(\lambda)} \right)^2.$$

Regression splines There are many different types of splines and different directions to go in. For example, one disadvantage of natural cubic smoothing splines is that we have a “parameter” of dimension n to estimate (namely, the vector \mathbf{g}). We can also reduce the number of knots to ξ_1, \dots, ξ_K and locate ξ_k at the $\left(\frac{k+1}{K+2}\right)$ -th sample quantile of x_1, \dots, x_n . Splines of order p can then be expanded in the *truncated power series basis*

$$1, x, x^2, \dots, x^p, (x - \xi_1)_+^p, \dots, (x - \xi_K)_+^p.$$

Therefore we can minimise the residual error over all polynomials in the span of this basis, which gives a parameter in \mathbb{R}^{p+1+K} to estimate using least squares. The solution is called a *regression spline*. Here, K plays the role of the smoothing parameter.

4 Minimax lower bounds

We seek lower bounds on the worst-case risk of any procedure, which provide a ‘benchmark’ against which we can measure the performance of a proposed method.

4.1 Reduction to testing

We will assume our parameter space (Θ, d) is a metric space or a semi-metric space (where we don’t require that $d(\vartheta, \vartheta') = 0 \implies \vartheta = \vartheta'$). We denote our collections of distributions depending on our parameters by $\{P_\vartheta \mid \vartheta \in \Theta\}$, which are probability measures on some measurable space $(\mathcal{X}, \mathcal{A})$.

Now, we let (Ω, \mathcal{F}) be any measurable space with a collection of probability measures $\{\mathbb{P}_\vartheta \mid \vartheta \in \Theta\}$ and a measurable function $X: \Omega \rightarrow \mathcal{X}$ so that $X \sim P_\vartheta$ on $(\Omega, \mathcal{F}, \mathbb{P}_\vartheta)$. Let $\hat{\Theta}$ denote the set of possible estimators for ϑ , i.e., all measurable functions $\mathcal{X} \rightarrow \Theta$.

Now, suppose we wish to estimate ϑ with a loss function of the form

$$L(\vartheta', \vartheta) = g(d(\vartheta', \vartheta)) \quad g \text{ increasing, } \vartheta', \vartheta \in \Theta.$$

We then define the *minimax risk* as

$$\mathcal{M} := \inf_{\hat{\vartheta} \in \hat{\Theta}} \sup_{\vartheta \in \Theta} \mathbb{E}_\vartheta L(\hat{\vartheta}(X), \vartheta),$$

i.e., the lowest worst-case estimated loss of any possible estimator $\hat{\vartheta}$.

Example 4.1. Suppose we are trying to estimate the mean of a normally distributed random variable with variance 1, and we have a sample (X_1, \dots, X_n) . Then $\Theta = \mathbb{R}$ with the Euclidian distance, $(\mathcal{X}, \mathcal{A}) = (\mathbb{R}^n, \mathcal{B}\mathbb{R}^n)$, and $P_\vartheta(F) = \int_F f_\vartheta(x_1) \cdots f_\vartheta(x_n) d\lambda^n(x)$, where f_ϑ is the density function of a $N(\vartheta, 1)$ distribution.

Now, let $X_\vartheta: (\Omega, \mathcal{F}) \rightarrow (\mathbb{R}^n, \mathcal{B}\mathbb{R}^n)$ be a $N_n(\vartheta \mathbf{1}, I_n)$ distributed random variable. In this case, \mathbb{P}_ϑ is a probability measure on Ω such that $P_\vartheta(F) = \mathbb{P}_\vartheta(X \in F)$.

Let $\hat{\Theta}$ denote the set of all estimators of ϑ , which are functions $\mathbb{R}^n \rightarrow \mathbb{R}$. Our loss function could simply be $L(\vartheta', \vartheta) = |\vartheta' - \vartheta|$ or $L(\vartheta', \vartheta) = (\vartheta' - \vartheta)^2$.

For $M \in \mathbb{N}$, let $[M] := \{1, \dots, M\}$, and let $\hat{\mathcal{T}}$ denote the set of measurable functions $\mathcal{X} \rightarrow [M]$. Given any $\vartheta_1, \dots, \vartheta_M \in \Theta$ and $\hat{\vartheta} \in \hat{\Theta}$, we can define $T_{\hat{\vartheta}} \in \hat{\mathcal{T}}$ by

$$T_{\hat{\vartheta}}(x) := \arg \min_{j \in [M]} d(\hat{\vartheta}(x), \vartheta_j),$$

where we pick the smallest j in case of a tie. Intuitively, we are simply approximating $\hat{\vartheta}$ by the closest ϑ_j . Now, we will lower-bound the minimax risk by an expression that only depends on estimators in $\hat{\mathcal{T}}$.

Writing $\eta = \frac{1}{2} \min_{j,k} d(\vartheta_j, \vartheta_k)$, we can lower-bound the worst-case loss of any fixed estimator $\hat{\vartheta}$ by

$$\begin{aligned} \sup_{\vartheta \in \Theta} \mathbb{E}_\vartheta L(\hat{\vartheta}(X), \vartheta) &\geq \max_{j \in [M]} \mathbb{E}_{\vartheta_j} g(d(\hat{\vartheta}, \vartheta_j)) \\ &= \max_{j \in [M]} \mathbb{E}_{\vartheta_j} \left\{ g(d(\hat{\vartheta}, \vartheta_j)) \mathbb{1}_{T_{\hat{\vartheta}} \neq j} \right\} \\ &\stackrel{\star}{\geq} g(\eta) \max_{j \in [M]} \mathbb{E}_{\vartheta_j} \mathbb{1}_{T_{\hat{\vartheta}} \neq j} \\ &= g(\eta) \max_{j \in [M]} P_{\vartheta_j}(T_{\hat{\vartheta}} \neq j), \end{aligned}$$

where \star holds because if $T_{\hat{\vartheta}} \neq j$, then $d(\hat{\vartheta}(x), \vartheta_j) \geq \eta$.

We therefore have

$$\begin{aligned}\mathcal{M} &\geq g(\eta) \inf_{\vartheta \in \hat{\Theta}} \max_{j \in [M]} P_{\vartheta_j}(T_{\hat{\vartheta}} \neq j) \geq g(\eta) \inf_{T \in \hat{\mathcal{T}}} \max_{j \in [M]} P_{\vartheta_j}(T \neq j) \\ &= g(\eta) \left\{ 1 - \sup_{T \in \hat{\mathcal{T}}} \min_{j \in [M]} P_{\vartheta_j}(T = j) \right\} \geq g(\eta) \left\{ 1 - \sup_{T \in \hat{\mathcal{T}}} \frac{1}{M} \sum_{j=1}^M P_{\vartheta_j}(T = j) \right\}.\end{aligned}$$

Therefore, we have now reduced the problem of lower-bounding \mathcal{M} to the problem of upper-bounding $\sup_{T \in \hat{\mathcal{T}}} \frac{1}{M} \sum_{j=1}^M P_{\vartheta_j}(T = j)$, which is a testing problem.

4.2 Divergences

Definition 4.2. Let μ, ν be measures on $(\mathcal{X}, \mathcal{A})$. We say that μ is *absolutely continuous* w.r.t. ν , notation $\mu \ll \nu$, if

$$\nu(A) = 0 \implies \mu(A) = 0.$$

We say that μ, ν are *mutually singular*, notation $\mu \perp \nu$, if there exists $A \in \mathcal{A}$ such that $\mu(A) = 0$ and $\nu(A^c) = 0$.

Note that mutual singularity means that μ “lives on” A^c , while ν “lives on” A .

Theorem 4.3 (Lebesgue). *If μ, ν are σ -finite measures on $(\mathcal{X}, \mathcal{A})$, then there exists measures μ_{ac} and μ_{sing} on $(\mathcal{X}, \mathcal{A})$ such that μ can be decomposed as $\mu = \mu_{\text{ac}} + \mu_{\text{sing}}$, where $\mu_{\text{ac}} \ll \nu$ and $\mu_{\text{sing}} \perp \nu$. Furthermore, this decomposition is unique.*